



Tier 1 Support Automation

Feasibility Overview

Prepared for Bulletproof, a GLI Company

Prepared by Bonito AI

May 2026

Confidential

The Problem

Bulletproof operates ~50 Tier 1 support agents across 300+ client sites in gaming, government, finance, healthcare, and manufacturing. These agents cost \$55–70K/year each — \$2.75M–\$3.5M in total annual labor.

Tier 1 handles the front line: high-volume, repetitive, well-documented work. The knowledge to resolve most tickets already exists in runbooks and SOPs — it's just trapped in documents that agents have to search manually.

~50

Tier 1 Agents

\$55–70K/yr each

\$3.5M

Annual Labor Cost

Tier 1 support alone

300+

Client Environments

Gaming, government, enterprise

This creates four compounding problems:

- **Slow response times** — agents read through docs while the user waits
- **Inconsistent quality** — resolution depends on which agent picks up the ticket
- **Shift-dependent coverage** — no 24/7 without overtime or offshoring
- **High turnover cost** — every new hire needs weeks of training on client-specific procedures

The work is well-structured. That's what makes it automatable. AI agents backed by client-specific knowledge bases can resolve the same tickets — faster, more consistently, and around the clock.



What the AI Handles

Bonito's agent framework deploys knowledge-backed AI agents that resolve common Tier 1 tickets autonomously — and intelligently escalate everything else with full context.

TICKET CATEGORY	% OF VOLUME	AI RESOLUTION
Password resets & account lockouts	~30%	Fully automated — KB lookup, identity verification, self-service routing
Connectivity & VPN issues	~20%	Automated troubleshooting; infrastructure issues escalated immediately
Software installation & access	~15%	Fully automated — approved software checks, provisioning tickets
"How do I..." questions	~15%	Fully automated — step-by-step from client KB
Hardware issues	~10%	AI handles triage & data collection; humans do physical work
Ticket triage & routing	~10%	Fully automated — classify, enrich, route to correct team

What Bulletproof Retains Control Of

- Which AI models the agents can use (per client, per environment)
- Budget caps per agent, per client — hard stops, not just alerts
- Human-in-the-loop approval gates on sensitive actions
- Per-client Knowledge Base isolation — government data never touches gaming inference
- Complete audit trail of every AI interaction, model call, and tool use

Escalation — Not Guessing

Every agent has confidence thresholds and hard rules. When a ticket exceeds scope, the agent escalates with the full conversation, KB articles referenced, steps attempted, and a suggested resolution. The human picks up mid-stream, not from scratch.

Compliance: Full audit trails for government municipality clients. Agent interactions logged and auditable for gaming clients (NIGC MICS, GLI). SOC 2 aligned controls across the board.



Cost Model

Current State

ITEM	VALUE
Tier 1 agents	~50
Avg salary + benefits	\$55,000–\$70,000/yr
Total annual Tier 1 labor	\$2,750,000–\$3,500,000
Avg tickets/month (estimated)	8,000–15,000

Projected State (Month 6+)

ITEM	VALUE
AI-deflected tickets	40–60%
Remaining human agents needed	20–25
Agents redeployed or reduced	25–30

Net Savings

Bulletproof uses their own cloud providers for inference — Bonito is the control plane, not the compute. LLM token costs are not a Bonito line item.

SCENARIO	CONSERVATIVE (40%)	MODERATE (50%)	AGGRESSIVE (60%)
Agents reduced	20	25	30
Labor saved/yr	\$1.1M–\$1.4M	\$1.4M–\$1.8M	\$1.7M–\$2.1M
Net savings/yr	\$920K–\$1.2M	\$1.2M–\$1.6M	\$1.5M–\$1.9M
ROI	6–7x	7–9x	8–11x

~\$1M+

Annual Savings

Even at conservative 40% deflection

6–11x

Return on Investment

Year 1 including implementation

\$0

Inference Cost to Bonito

Runs on Bulletproof's own cloud

Rollout Timeline

Enterprise deployment across 300+ client environments with compliance requirements. Five phases, fully production in 6–8 months.

PHASE 1 — WEEKS 1–4

Discovery & KB Ingestion

Runbooks and SOPs ingested for pilot clients. Coverage gaps identified. Escalation rules defined per client.

PHASE 2 — WEEKS 5–8

Agent Build & Integration

Agents configured per client group. ITSM integration (ServiceNow, ConnectWise). Internal testing with synthetic tickets.

PHASE 3 — WEEKS 9–12

Shadow Mode Pilot

Agents deployed in shadow mode on 3–5 clients. Human agents see AI suggestions and choose to use, modify, or ignore. Zero risk — humans still resolve every ticket.

PHASE 4 — WEEKS 13–16

Supervised Go-Live

Live on pilot clients starting with low-risk categories. Human spot-checks on 20% of resolutions. Gradually expand as accuracy validates.

PHASE 5 — MONTHS 5–8

Scale to Full Client Base

Roll out in waves (20–30 clients per wave). Headcount reduction begins. Resolved tickets feed back into KB for continuous improvement.

Why Bonito (vs Build It Yourself)

- **Multi-provider failover** — if one AI provider goes down at 2 AM, Bonito reroutes automatically. Support agents are 24/7; provider uptime is not.
- **Per-client tenant isolation** — architecture-level, not prompt-level. Government data stays separated from gaming inference.

- **Compliance-ready** — audit trails, HIPAA/GDPR policy enforcement, SOC 2 aligned controls. Built in, not bolted on.
- **Cost visibility & hard stops** — exact cost per agent, per client, per ticket. Budget caps that halt spend, not just alert.
- **No infrastructure to maintain** — Bonito handles model updates, scaling, and uptime. Bulletproof configures agents and manages KBs.

Next Step

Happy to walk through the details or discuss what a scoped pilot would look like — we can tailor it to however your team wants to approach the build.

Bonito AI

shabari@bonito.ai

<https://getbonito.com>